

OSSAMA ELHAKKI

ML Infrastructure & Deep Learning Engineer — Training · Optimization · Deployment

+212 6 73 10 52 17 · ossamaelhakki@gmail.com · Casablanca, Morocco · ossamaelhakki.com
linkedin.com/in/elhakki-ossama-985b64190 · github.com/ELHAKKI-OSSAMA · kaggle.com/ossamaelhakk

ABOUT ME

Deep Learning engineer (Master of Science in Distributed Systems & AI, Applied Mathematics background) focused on **training, validating and optimizing deep neural networks** and on the tooling and infrastructure around ML workflows. Experienced in **model optimization and deployment** (ONNX export, CUDA, inference serving) and in setting up reproducible training/validation pipelines. Strong in **Python** and deep-learning frameworks (PyTorch, TensorFlow), with solid maths fundamentals and good English communication. Eager to apply DL optimization to resource-constrained targets.

EXPERIENCE

Ofoto — Fine Art Digital Studio

Deep Learning & ML Infrastructure Engineer | Feb 2025 – Oct 2025 | Casablanca

- **DL model deployment & optimization** — served deep-learning models (Stable Diffusion, InstantID) in production, optimizing throughput and resource usage — -35% latency, 500+ concurrent requests.
- **Inference infrastructure** — built serving/queueing infrastructure (Redis, GPU-backed inference) and applied CUDA/ONNX-oriented optimizations for efficient runtime.
- **ML tooling & workflows** — set up reproducible deployment pipelines with Docker and CI/CD (GitHub), plus monitoring (Prometheus/Grafana).
- **Cross-team collaboration** — coordinated across frontend, backend and infrastructure, with clear technical documentation.

FSTM Research Laboratory

Deep Learning Research Intern — Computer Vision | Feb 2023 – Jul 2023 | Mohammedia

- Trained and validated a real-time CNN-based face-recognition pipeline (OpenCV), running inference against a live database.

SELECTED DEEP LEARNING WORK

- **Model training & validation** — trained and benchmarked CNNs, RNN/LSTM, transformers and generative models (DCGAN) in PyTorch/TensorFlow with rigorous validation.
- **Optimization & export** — ONNX export and inference acceleration (CUDA); hyperparameter optimization with Optuna; walk-forward benchmarking of 19 models.
- **Serving & MLOps** — model serving with KServe, experiment tracking with MLflow, monitoring with Prometheus/Grafana.

EDUCATION

ENSET Mohammedia — M.Sc., Distributed Systems & AI	<i>Bien</i>
FST Mohammedia — B.Sc., Applied Mathematics	<i>Assez Bien</i>

SKILLS

Languages	Python, C++ (foundations), Java, SQL
Deep Learning	PyTorch, TensorFlow, Keras; CNN, RNN/LSTM, Transformers, U-Net, EfficientNet, DCGAN
DL Optimization	ONNX export, CUDA, inference acceleration, quantization & pruning concepts, Optuna
ML Infra / MLOps	KServe, MLflow, Docker, Kubernetes, CI/CD, Prometheus + Grafana, GPU workflows
Training & Validation	Reproducible pipelines, cross-validation, benchmarking, evaluation metrics
Maths	Applied mathematics, linear algebra, probability & statistics

LANGUAGES

Arabic — Native · English — B1/B2 · French — B1/B2